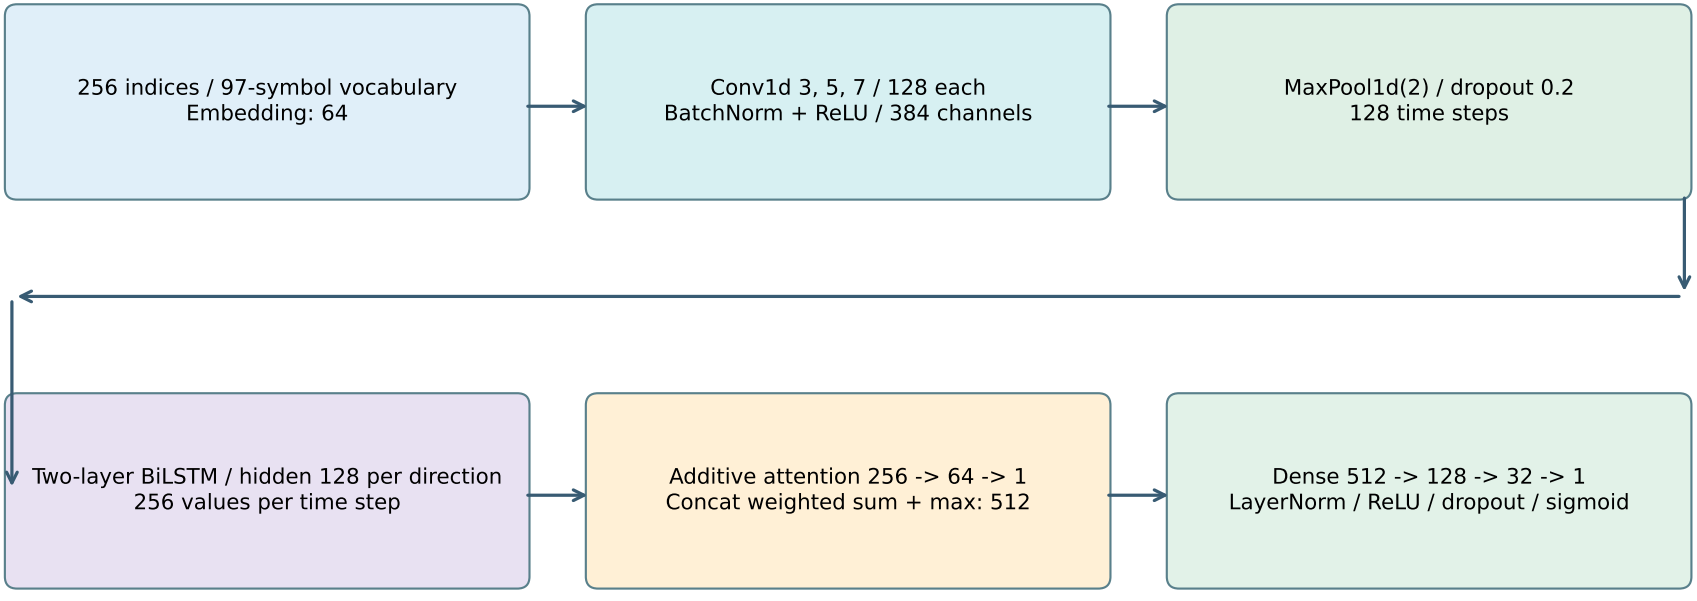
256 indices / 97-symbol vocabulary
Embedding: 64
Conv1d 3, 5, 7 / 128 each
BatchNorm + ReLU / 384 channels
MaxPool1d(2) / dropout 0.2
128 time steps
Two-layer BiLSTM / hidden 128 per direction
256 values per time step
Additive attention 256 -> 64 -> 1
Concat weighted sum + max: 512
Dense 512 -> 128 -> 32 -> 1
LayerNorm / ReLU / dropout / sigmoid